Qwen

From Theory to Telemetry: A Blueprint for a Cryptographically-Auditable Horror Asset Factory

Foundational Infrastructure Spine

The establishment of a foundational infrastructure spine is the primary and non-negotiable first step in creating a coherent and scalable VM constellation for HorrorPlace . This approach prioritizes the creation of a rigid, shared technical standard across all participating repositories to prevent divergence and ensure that all subsequent content experiments are built upon a consistent and machine-verifiable foundation . The core risk of developing content in parallel without a firm infrastructure backbone is the emergence of incompatible data models and workflows, rendering assets irreconcilable and the entire system unmanageable . Therefore, the initial phase of development must focus exclusively on defining and implementing the shared rules of engagement that will govern every aspect of the asset creation and validation process. This involves locking down a suite of versioned JSON schemas, standardizing CI/CD pipelines through reusable templates, and establishing clear documentation for new vault onboarding. The objective is to create a "spine" of infrastructure that all repositories—both existing and future—will adhere to, forming the structural integrity of the entire circulatory system for horror.

A critical component of this spine is the definition of a comprehensive set of shared, versioned JSON schemas. These schemas act as the single source of truth for all data structures used within the ecosystem, ensuring that concepts like "invariants," "entertainment metrics," and "agent artifacts" have a universally understood and strictly enforced structure . The provided context highlights the necessity of several key schema types. First, the `schemas/invariants_v1.json` defines the bundle of numeric properties that encode the geo-historical horror profile of any spatial unit in the game world, including CIC (Catastrophic Imprint Coefficient), AOS (Archival Opacity Score), SHCI (Spectral-History Coupling Index), and other derived values like DET (Dread Exposure Threshold) and LSG (Liminal Stress Gradient) . Second, the `schemas/entertainment_metrics_v1.json` provides the definitions and validation rules for machine-checkable metrics that quantify player experience, such as UEC (Uncertainty Engagement Coefficient), EMD (Evidential Mystery Density), STCI (Safe-Threat Contrast

Index), CDL (Cognitive Dissonance Load), and ARR (Ambiguous Resolution Ratio) . Third, a family of declarative layer schemas—including `schema://Horror.Place/core/schemas/stylecontract_v1`, `schema://Horror.Place/core/schemas/event_contract_v1`, and `schema://Horror.Place/core/schemas/persona_contract_v1`—defines the rules and constraints for creating new content, acting as cross-tier firewalls by specifying what evidence is permissible and explicitly forbidding explicit depictions . Finally, the infrastructure requires schemas for artifacts and provenance, namely `schema://Horror.Place/core/schemas/agent_artifact_v1` and `schema://Horor.Place/core/schemas/dead_ledger_entry_v1`, which provide the structure for cryptographically signing AI personas and events and tracking their performance and lineage . The existence of these schemas, mirrored from the public `Horror.Place` core repository, is the prerequisite for any meaningful development in the private vaults and labs .

Complementing the schemas is the implementation of standardized Continuous Integration/Continuous Deployment (CI/CD) pipelines. The provided materials offer production-ready CircleCI 2.1 template files designed to enforce the specific workflow of the HorrorPlace ecosystem . Two distinct templates are necessary to cover the different roles within the constellation. The first is a general-purpose template for orchestrator-type services, `ci-templates/circleci-orchestrator-template.yml`, which establishes a baseline of quality control for any service consuming or managing artifacts. This template enforces a strict pipeline of linting (using `ruff`), unit testing with a minimum coverage threshold of 70%, integration tests, and a final content leak scan that specifically looks for raw asset URLs and private keys . The second, more specialized template, `ci-templates/circleci-vault-validate-template.yml`, is tailored for the vault repositories (`Atrocity-Seeds`, `Black-Archivum`, etc.) and automates the entire lifecycle of a seed artifact . This pipeline begins by validating all seed files against their declared public schemas, proceeds to compute canonical SHA-256 hashes for integrity verification, calls a hardware-backed Key Management Service (KMS) to sign the artifacts, produces a structured `artifact-descriptor.json`, and concludes by dispatching a `repository_dispatch` event back to the orchestrator repository . The use of these pre-built, merge-ready templates ensures that every repository adheres to the same rigorous standards for code quality, security, and artifact handling, preventing subtle but critical differences in how builds are processed . The integration guide further specifies that these templates rely on CircleCI contexts for sensitive information like KMS endpoints and GITHUB_TOKENs, which are never stored directly in the repository, reinforcing the security posture from the outset .

Finally, the foundational infrastructure is completed by establishing clear documentation and procedures for onboarding new repositories. The `ops/kms_hsm_integration.md`

document serves as a master guide, but a dedicated checklist is needed for new vaults . This checklist outlines a systematic process: request a signing key from the security team, receive the `KMS_KEY_ID` and `KMS_ENDPOINT`, configure the necessary CircleCI context, add the corresponding public key reference to the orchestrator's `config/trusted_keys.json` file via a pull request, run a test pipeline to validate the setup, and only then enable the production pipeline on the new repository . The `trusted_keys.json` file itself has a strict format, containing the `key_id`, `algorithm`, `public_key` material, and metadata about its status and associated vault repositories . By codifying this process, the organization ensures that every new participant in the constellation is brought up to speed consistently and securely, maintaining the overall integrity of the system. The dependency graph provided in the materials clarifies how these components interrelate: the vault repositories consume the CI templates and refer to the KMS/HSM integration guide; the orchestrator consumes the orchestrator template and also refers to the KMS guide for verification; and both templates depend on the principles laid out in the comprehensive `kms_hsm_integration.md` document . This interconnectedness underscores that the infrastructure is not a collection of isolated parts but a cohesive whole where each piece relies on the others to function correctly. Together, the schemas, CI/CD templates, and onboarding procedures form the essential spine upon which all creative work will be built, providing the stability and predictability required for a large-scale, distributed project focused on evolvable horror assets.

Standardized Runtime and Experimentation Harness

While the foundational infrastructure provides the structural backbone, the standardized runtime and experimentation harness serve as the operational engine that drives the VM constellation. This component is crucial for translating abstract schemas and policies into tangible, testable implementations while ensuring that all outputs are comparable and reproducible. The user's directive to implement a reference Lua "Horror Runtime Core" with narrow APIs is a central tenet of this strategy . This runtime acts as a shared library of functions that all underground repositories must use, effectively creating a sandboxed environment where logic is executed consistently regardless of the underlying engine (e.g., Godot, a custom Lua VM) . The concept of a "narrow API" is paramount; it means that instead of allowing direct access to raw tables of invariants or metrics, all queries and interactions must go through a curated set of functions. For instance, a persona would call `H.CIC(region_id)` and `H.AOS(region_id)` rather than accessing a global table of invariants directly . This disciplined approach prevents ad-hoc logic that

could bypass the intended constraints and makes static analysis feasible, allowing tools to automatically verify that all code properly interacts with the horror engine's state . The documented reference module includes APIs for querying history (`H.`), reading entertainment metrics (`Metrics.`), receiving BCI inputs (`BCI.`), issuing directives (`Director.`), manipulating story spines (`Story.`), enforcing policies (`Policy.`), and controlling the Archivist personality (`A.`) . Enforcing the use of this shared runtime is a critical part of the CI/CD pipeline, likely through linting scripts that flag any direct access to forbidden data structures.

To facilitate parallel content experiments, a clearly documented and repeatable "experiment harness" is required. This harness is designed to allow researchers in each lab to generate "seed telemetry" efficiently and cheaply, using either synthetic data or a minimal simulation environment . The recommended approach involves standing up a canonical "Horror$Place Testbed VM" . This would be a small, self-contained project (e.g., in Godot or Love2D) that includes the shared Lua runtime, simple invariant-tagged regions, and basic visual corridors . All labs would use this identical testbed for initial experiments, ensuring that telemetry generated by Neural-Resonance-Lab, Redacted-Chronicles, and Process-Gods-Research is directly comparable because it originates from the same controlled environment . To further accelerate development, especially for privacy-safe large-batch runs, a Headless Simulation Mode should be developed . This mode would be a command-line interface (CLI) tool written in Lua that runs scripted pseudo-players through the generated environments, triggering events and persona interactions without requiring a full graphical render, thus emitting telemetry very quickly . This allows for rapid "seed burn-in" and extensive testing of persona behaviors under various conditions. Additionally, providing sandbox adapters for different input sources (e.g., one adapter for off-the-shelf EEG headbands, another for webcam-based emotion detection) would allow researchers to plug in different data streams without altering the core BCI schema, enhancing flexibility . The experiment harness must also define clear logging levels: raw data remains Tier-3 only, redacted logs are for Tier-2 consumption, and summarized telemetry is made available to Tier-1 . Each level would have a strict schema enforced by CI checks, ensuring data flows appropriately through the tiers without leaking sensitive information .

The role of each specific repository within this experimental framework is clearly defined by their conceptual purpose. The `Codebase-of-Death` acts as the primary implementer, tasked with taking a persona or event contract from Tier-1 and writing the Lua code that brings it to life, strictly using the narrow APIs . Its initial experiment would be to produce a small, fully compliant persona/event bundle that proves the entire loop from public contract to underground implementation works . The `Neural-Resonance-Lab` focuses on neuro-responsive horror, implementing BCI-driven adjustments to events

and defining the mapping between physiological signals (like a `fear_index`) and metric targets . Their first pass might involve wiring a prototype test in the shared testbed to apply synthetic BCI data and adjust a haptic pattern based on predefined rules . The `Redacted-Chronicles` lab's job is to act as a normalizer, taking the raw telemetry from playtests and compressing it into anonymized, signed summaries suitable for consumption by Tier-2 and Tier-1 systems . They would start by defining the three schema levels (raw, redacted, summary) and implementing pipelines to convert between them, running a toy experiment to verify the process . `Process-Gods-Research` specializes in extreme SHCI personas, and their initial task is to create a persona contract and a stub implementation of a neutral system (e.g., a water-distribution network) that is wired to respond to invariants but contains no explicit horror content, thus validating the core logic . Finally, `Liminal-Continuum` validates the entire pipeline from the other side; its first experiment is to accept a single signed agent artifact from one of the other labs and run it through its marketplace ranking flow, confirming that the acceptance criteria and validation scripts are working . This parallel, targeted approach ensures that each part of the system is validated in isolation before being integrated into the larger constellation, mitigating risk and building confidence in the overall architecture.

Cryptographic Fabric and Asset Validation

The cryptographic fabric is the bedrock of trust and integrity within the HorrorPlace VM constellation. It is the mechanism that ensures every piece of content moving through the asset factory—from a seed in a vault to a signed agent in the marketplace—is authentic, untampered-with, and compliant with the governing rules. The user's requirement for a system that enables a "seed-generating asset factory for horror assembly" hinges entirely on the robustness of this cryptographic enforcement, which moves beyond simple version control to active, verifiable governance of the creative process . The design philosophy is centered on hardware-backed signing and a multi-stage verification process that spans from the moment an artifact is created until it is consumed by the orchestrator. This system directly implements the "Rivers of Blood Charter doctrine" by making compliance with invariants, metrics, and implication-based rules a cryptographic property of the asset itself . The architecture is designed to be resilient, secure, and auditable, with clear controls over key lifecycle, disaster recovery, and operational security.

The core of the cryptographic system is the Key Management Service (KMS) / Hardware Security Module (HSM) integration, which is meticulously detailed in the `ops/kms_hsm_integration.md` document . The fundamental principle is absolute: private

keys must never leave the KMS/HSM boundary . No repository may store, transmit, or log private key material under any circumstances. All signing operations are performed remotely via a secure API call at the hardware boundary during the CI pipeline . This is enforced by a mandatory CI leak scan that runs on every pipeline execution to check for any instances of private key material in the code or logs . The Vault CI pipeline is responsible for generating a canonical hash of a seed file and then calling the KMS signing API with this hash. The KMS responds with a digital signature, which is then bundled with the seed file and its metadata into an `artifact-descriptor.json` . This descriptor, containing the signature, is what gets dispatched to the orchestrator. Upon receipt, the orchestrator's CI pipeline uses a separate KMS endpoint to call the verification API, passing the canonical hash and the signature from the descriptor. The HSM verifies the signature against the trusted public key and returns a simple boolean result: `valid: true` or `valid: false` . This two-way communication (sign and verify) forms a closed-loop integrity check that guarantees an artifact's provenance and authenticity before it can ever be processed further. The entire process is automated and gated by the CI/CD workflows, making manual bypass impossible.

Security and operational resilience are paramount, and the KMS/HSM guide addresses these concerns comprehensively. Security controls include short-lived authentication tokens for the CI pipeline (rotated per run, maximum lifetime of one hour), network-level access restrictions that limit the KMS API to known CI runner IP ranges, and IAM policies that grant keys the least privilege necessary, specifically scoping the `KMS_AUTH_TOKEN` to the `sign` operation only . Furthermore, all KMS operations are logged in a persistent audit trail, providing a record for review and compliance . The use of FIPS 140-2 Level 3 or higher certified HSMs ensures the physical and logical security of the key material itself . Disaster recovery scenarios are also planned for, including graceful failure if the KMS becomes unavailable (the pipeline simply fails, producing no unsigned artifacts), and a detailed procedure for emergency offline signing in the event of a prolonged provider outage . Key lifecycle management is another critical aspect, with a defined schedule for key rotation (every 90 days for standard keys), a multi-stage state model (Pending Activation, Active, Pre-Deprecated, Deprecated, Destroyed), and a minimum of two active keys at all times to support zero-downtime rotation and multi-party approval for destruction .

A sophisticated feature of the cryptographic fabric is the support for Multi-Party Signing, which is enabled for high-value artifacts . For Tier 1 artifacts, which are defined as those affecting public registries or changing core schemas, two signatures from different keys held by different operators are required . The process involves the vault CI signing with an operator key, and then the orchestrator, upon verifying that signature, requiring a second signature from a separate approver key before accepting the artifact and updating

the registry . This quorum-based system prevents a single point of failure or compromise and adds an additional layer of human oversight for critical changes. The configuration for this is managed via a `multi_party_signing` block in the orchestrator's config, which can be enabled or disabled and configured to require 2-of-N signatures for specific tiers of artifacts . This combination of hardware-backed signing, stringent security controls, robust disaster recovery planning, and multi-party approval creates a cryptographic fabric that is not merely an add-on but the central nervous system of the entire asset validation process. It transforms the asset factory from a simple pipeline into a secure, trustworthy, and auditable system for generating and distributing horror content.

Repository Roles and VM Constellation Workflow

The effectiveness of the HorrorPlace VM constellation depends on a well-defined division of labor, assigning specific roles and responsibilities to each of the designated repositories. This role-based model ensures specialization, prevents code duplication, and facilitates a smooth, predictable workflow for generating and validating horror assets. The VM constellation is not just a buzzword but a functional blueprint for a collaborative, multi-agent system where each node performs a distinct task in a tightly orchestrated sequence . The repositories are organized into a tiered architecture, with Tier 2 holding vaults of assets and Tier 3 housing specialized research labs, all interacting through GitHub-native orchestration . The workflow begins with seeds in Tier-2 vaults, progresses to tuning and validation in Tier-3 labs, and culminates in a ranked marketplace in another Tier-2 vault, from which preferred implementations are promoted back to the public Tier-1 core .

The specific roles for the highlighted repositories are highly specialized, reflecting their unique contributions to the asset factory.

Repository	Tier	Role(s) in Constellation	Primary Data/Logic
Codebase-of-Death	2	Breeder + Implementer VM	Implements personas and events as code and signed agent artifacts, strictly under persona/event contracts from Tier-1; primary home for evolved Archivist/Witness/Echo/Process/Threshold behaviors.
Neural-Resonance-Lab	3	Researcher VM	Ingests BCI/EDA/HRV data, defines BCI→metric mapping rules, haptic patterns, and neural resonance traits; produces aggregated, signed telemetry summaries back to Tier-2 and Tier-1.
Redacted-Chronicles	3	Researcher + Normalizer VM	Stores raw fMRI/BCI logs and detailed physiological traces; responsible for compressing them into anonymized metric-delta summaries suitable for Liminal-Continuum and Tier-1 consumption.
Process-Gods-Research	3	Breeder (extreme SHCI) VM	Specializes in Process-type entities (high FCF, low DET, high SHCI) that describe systems consuming without malice; explores existential dread under strict invariant/metric constraints.
Liminal-Continuum	2	Seeder + Marketplace VM	Implements agent artifact schemas and Dead-Ledger entries; ranks agents by invariant fidelity and entertainment metrics; defines agent exchange protocol and roles (Seeders, Breeders, Validators, Researchers).
Obscura-Nexus	2	Seeder (styles/events) VM	Hosts experimental style contracts (e.g., Rivers of Blood Spectral Fluid Chrome family), surprise event variants, and research-tier style extensions; all expressed via style contract DSL.

This division of labor creates a powerful feedback loop. The workflow for a new asset begins with a seed, perhaps an invariant-bound event contract from `Atrocity-Seeds` . `Codebase-of-Death` takes this contract and generates a conformant Lua implementation, which is then packaged as a signed agent artifact . This artifact is dispatched to the orchestrator, which in turn notifies the relevant labs. `Neural-Resonance-Lab` and `Redacted-Chronicles` would then run playtests on this agent in the shared testbed, correlating BCI data and other telemetry with the agent's performance . The resulting compressed telemetry summaries from `Redacted-Chronicles` and the rule sets from `Neural-Resonance-Lab` are fed back into the system . Meanwhile, `Obscura-Nexus` might provide stylistic variations for the event, creating different `style_contract` instances that can be applied to the base implementation . The `Process-Gods-Research` lab could generate alternative persona implementations that fulfill the same contract but exhibit different behavioral patterns under the constraint of high SHCI .

Once these experiments are complete, the results converge in `Liminal-Continuum`. This repository acts as the central hub and validator. It ingests all the signed agent artifacts and their associated telemetry summaries . Using the cryptographic and entertainment checks defined in the VM-constellation spec, it validates each agent's signature and invariant compliance . Only after passing these initial checks are the agents' entertainment efficacy metrics (e.g., metric lift, ARR, variance) considered . `Liminal-Continuum` then applies a ranking function, prioritizing agents with flawless

invariant compliance followed by positive metric lift and stable performance across multiple sessions . The outcome of this process is a ranked list of "preferred implementations" for each persona and event contract ID. These rankings are logged in a `Dead-Ledger` and published back to the Tier-1 `registry_dead_ledger.json` . This completes the loop, as the public `Horror.Place` core now has a verifiable, machine-trusted source for the best available implementations of its abstract contracts. The entire system operates as a closed, GitHub-native workflow driven by `repository_dispatch` events, where a change in one repository triggers a chain reaction of processing and validation in others, creating a truly circulatory system for evolving horror assets .

Success Metrics and Safety Guardrails

Defining a framework for success is as critical as building the infrastructure itself, as it provides the criteria for evaluating experiments and guiding the evolution of the horror assets. The approach proposed for HorrorPlace is a hybrid model that balances hard, non-negotiable constraints (guardrails) with softer, exploratory goals (targets) . This dual-layer system ensures that the creative process remains within safe, ethical, and technically sound boundaries while still allowing for the discovery of novel and effective horror experiences. The guardrails are absolute requirements for safety, integrity, and quality, enforced automatically by the CI/CD system. The soft targets, conversely, represent the desired direction of the creative exploration, encouraging innovation within a defined range rather than forcing adherence to a single, rigid outcome.

The hard constraints, or guardrails, are designed to prevent catastrophic failures and maintain the system's integrity. The first and most fundamental guardrail is **Invariant Fidelity**. Every persona and event implementation must rigorously obey the local invariants of the region in which it operates. Code must be linted and statically analyzed to ensure that all logic correctly reads invariants through the narrow `H.` API and does not introduce "free-floating" scares that are disconnected from the historical and spectral context of the location . This is a core tenet of the Rivers of Blood Charter's doctrine of implication-based horror . The second guardrail is **Cryptographic Hygiene**. The system mandates 100% signature validation for all agent artifacts entering the `Liminal-Continuum` marketplace. Any artifact that fails signature verification or contract validation must be automatically rejected and logged; there is no manual override . This ensures that only cryptographically verified and compliant code can ever be deployed. The third category of guardrails relates to **BCI Safety and Comfort Envelopes**. Given the use of neuro-responsive elements, explicit safety limits must be defined. These include

maximum sustained fear_index levels, cumulative exposure time above a certain threshold, and built-in downshift triggers that reduce event intensity or increase safe zones if these thresholds are exceeded . This protects participants from undue distress and ensures the system remains a tool for "fear-as-fun" rather than genuine anxiety . Finally, all content must adhere to a strict **Entitlement and Implication Model**. Even within the private research labs, content cannot be explicit; it must remain within the bounds of implication, respecting tiered intensity models controlled by feature flags .

Within these firm guardrails, the framework encourages exploration guided by soft targets and metric bands. The goal is not to hit a single magic number but to achieve a desirable *range* of outcomes that contribute to an engaging player experience. For adult research runs, success is measured against **Metric Bands** for entertainment quality. For example, a successful session might aim for a UEC (Uncertainty Engagement Coefficient) between 0.55 and 0.8, an ARR (Ambiguous Resolution Ratio) greater than or equal to 0.7, and a CDL (Cognitive Dissonance Load) that is kept at a "moderate" level appropriate for the region's theme, avoiding either collapse into pure confusion or premature resolution . The exact pacing curves and shapes of metric trajectories are not fixed targets but rather hypotheses to be tested and refined based on telemetry data . Similarly, the evolution of AI personas like the Archivist or Process is treated as an evolutionary process. Success is defined as an ability to improve UEC and ARR under the given invariant and metric constraints across a sufficient number of runs (e.g., N=10), rather than reaching a predetermined score . This allows the evolutionary strategies to discover emergent and surprisingly effective behaviors that might not have been conceived initially, fostering true innovation within the system's rules .

This hybrid model is operationalized through the experiment design and telemetry pipelines. Labs are encouraged to run hypothesis-driven experiments, such as "10-run microstudies" for each new agent, which track per-phase metrics and BCI readings, computing averages and variance . An agent is only considered for promotion to the marketplace if its variance is stable and it respects all guardrails . Another method is "A/B style trials" in `Obscura-Nexus`, where two different style contracts are bound to the same event to compare their resulting metric distributions and refine the style envelopes accordingly . The telemetry itself is structured to reflect this philosophy. While raw data is confined to Tier-3, the redacted and summary telemetry produced by `Redacted-Chronicles` and `Neural-Resonance-Lab` is rich with contextual information, including the specific invariant band the session was run in, the agent or event ID, and the computed metric deltas . This allows `Liminal-Continuum` to rank agents not just on their absolute performance but on how well they perform *relative to the constraints imposed by their environment*. For example, an agent that achieves a high ARR in a high-AOS region might be ranked higher than one that achieves the same ARR in a low-AOS

region, as the former demonstrates a better understanding of its narrative context. By combining unyielding guardrails with an exploratory, data-driven approach to success, the framework ensures that the HorrorPlace asset factory is both a safe and a creative engine for generating novel, effective, and explainably horrifying experiences.

Phased Rollout Plan and Synthesis

The culmination of this research framework is a concrete, phased rollout plan that translates the architectural vision into an actionable roadmap for the research teams. This plan prioritizes the sequential delivery of capabilities, starting with the foundational infrastructure and progressively layering on content experiments and advanced features. The overarching goal of this phased approach is to build a functional, end-to-end circulatory system for horror assets, where a seed can be generated in a vault, tested in a lab, validated in a marketplace, and its improved implementation returned to the core rules engine . The plan is structured around a series of milestones, each representing a significant step toward a fully operational VM constellation. This synthesis of infrastructure, roles, and metrics provides a clear path forward, ensuring that the development across the five repositories is aligned and incremental progress is continuously validated.

The initial phase, spanning approximately 12 weeks, is dedicated to establishing the foundational infrastructure spine. During this period, the primary focus is on lock-down and implementation, with minimal content creation. The first two weeks are allocated to defining and sharing the core schemas and the initial CI/CD templates. Repositories like `Neural-Resonance-Lab` and `Redacted-Chronicles` would immediately adopt a pilot specification for BCI payloads and telemetry summaries, even if it is based on synthetic data, to align their logging formats from day one . Concurrently, `Codebase-of-Death` and `Liminal-Continuum` would implement their respective agent artifact and dead ledger schemas, setting the stage for cryptographic validation . Weeks 3-4 would see the implementation of the CI/CD pipelines in each repository, integrating the CircleCI templates, setting up the necessary CircleCI contexts for secrets, and adding initial linting and schema validation checks to the build process . The next four weeks (Weeks 5-8) would focus on refining the narrow Lua APIs and the shared testbed environment. `Codebase-of-Death` would write its first minimal, compliant agent implementation (e.g., a Process persona stub), while `Neural-Resonance-Lab` develops its BCI adapter and `Redacted-Chronicles` builds its telemetry ingestion and redaction pipelines . The final four weeks of this phase (Weeks 9-12) are for integration

and validation. `Liminal-Continuum` would implement its artifact validation and ranking scripts, and the team would conduct a dry-run, feeding a single, manually created signed agent artifact through the entire pipeline to ensure all components—validation, dispatch, and ranking—function as expected .

Following the successful completion of Phase 1, the project enters a parallel development and validation phase. With the infrastructure in place, the focus shifts to introducing "thin, very targeted content experiments" in each repository to keep the loop real and gather actual data . During this phase, `Neural-Resonance-Lab` would move from synthetic data to running small-scale experiments with real or simulated BCI inputs, implementing the first BCI-to-metric rules and haptic pattern mappings . `Redacted-Chronicles` would begin running local play sessions, compressing the resulting telemetry into the formal summary schema and publishing it to its designated storage . `Obscura-Nexus` would expand its work on the Spectral Fluid Chrome style family, authoring new style contracts and event variants bound to them . `Process-Gods-Research` would evolve its Process persona contracts, testing them in controlled environments and adjusting their metric profiles based on the logged outcomes . Throughout this phase, `Liminal-Continuum` becomes the active consumer, ingesting the telemetry and agent artifacts from the other labs, running them through its validation and ranking system, and populating its internal registry with preliminary performance data. This phase is characterized by rapid iteration and feedback, with each lab's findings informing the others and driving refinements to the schemas and experimental designs.

The final phase transitions the system from a blueprint to a functioning, continuous evolution engine. Once at least one complete Surprise.Event! contract has run end-to-end—from its Tier-1 specification, through its implementation and testing in the underground repos, the generation of telemetry, and its potential promotion back into the Tier-1 registries—the VM constellation is considered operational . The long-term objective is to leverage the Seeders, Breeders, Validators, and Researchers role model to iteratively evolve agents based on ongoing metric and BCI feedback . Work would focus on scaling the orchestration, potentially moving beyond simple `repository_dispatch` calls to more robust messaging systems if bottlenecks arise. The Dead-Ledger Network would be fully implemented to provide transparent traceability of agent lifecycles . Most importantly, the system would begin to automate the refinement of its own rules. Based on aggregated telemetry, the system could automatically propose updates to persona or style contracts, suggesting new metric targets or invariant constraints that better capture what produces effective horror. This creates a virtuous cycle where the engine learns from its own output, leading to increasingly sophisticated and nuanced horror assets. This final phase represents the achievement of the research goal: a self-sustaining,

cryptographically-auditable asset factory for horror assembly, capable of generating "scare-ready agents" that are not just scary, but explainably and repeatably so .

Reference

1. Artificial Intelligence and Games - Springer Nature https://link.springer.com/content/pdf/10.1007/978-3-031-83347-2.pdf
2. The Architecture of Open Source Applications, Volume II | PDF - Scribd https://www.scribd.com/document/199572200/The-Architecture-of-Open-Source-Applications-Volume-II
3. 333333 23135851162 the 13151942776 of 12997637966 https://www.cs.princeton.edu/courses/archive/spring18/cos226/assignments/autocomplete/testing/words-333333.txt
4. hw3_stats_google_1gram.txt - CMU School of Computer Science https://www.cs.cmu.edu/~roni/11661/2017_fall_assignments/hw3_stats_google_1gram.txt
5. Festschrift https://www.microsoft.com/en-us/research/wp-content/uploads/2012/07/Rick-Rashids-Festschrift-7.5.12.pdf
6. Emerging Trends and Technologies on Intelligent Systems https://link.springer.com/content/pdf/10.1007/978-981-97-5703-9.pdf
7. Extreme Java - Concurrency and Performance For Java 8 Course ... https://www.scribd.com/document/373422261/Extreme-Java-Concurrency-and-Performance-for-Java-8-Course-Preparation-Reading
8. 333333 23135851162 the 13151942776 of 12997637966 ftp://ftp.cs.princeton.edu/pub/cs226/autocomplete/words-333333.txt
9. Extended Reality - Springer Link https://link.springer.com/content/pdf/10.1007/978-3-031-97763-3.pdf
10. Google 1-gram - CMU School of Computer Science https://www.cs.cmu.edu/~roni/11761/2017_fall_assignments/hw3_stats_google_1gram.txt
11. [PDF] Technical Communication for Environmental Action - ResearchGate https://www.researchgate.net/profile/Daniel-Richards-6/publication/368502558_Technical_Communication_for_Environmental_Action/links/662260fd43f8df018d1a00ad/Technical-Communication-for-Environmental-Action.pdf

12. [PDF] The Competitiveness of Global Port-Cities: Synthesis Report - OECD https://www.oecd.org/content/dam/oecd/en/publications/reports/2013/09/the-competitiveness-of-global-port-cities-synthesis-report_g17a2367/5k40hdhp6t8s-en.pdf
13. Multicriteria decision analysis in water resources management https://unesdoc.unesco.org/ark:/48223/pf0000097402
14. Sensors, Volume 26, Issue 4 (February-2 2026) – 321 articles - MDPI https://www.mdpi.com/1424-8220/26/4
15. Debian -- Prospective packages https://www.debian.org/devel/wnpp/prospective
16. Social Cognition in Human-Robot Interaction - HRI - ResearchGate https://www.researchgate.net/publication/327270113_Social_Cognition_in_Human-Robot_Interaction_Putting_the_'H'_back_in_'HRI'
17. Sensors, Volume 25, Issue 15 (August-1 2025) – 353 articles https://www.mdpi.com/1424-8220/25/15
18. (PDF) Decision Intelligence Frameworks for Sustainable Fisheries ... https://www.researchgate.net/publication/401954707_Decision_Intelligence_Frameworks_for_Sustainable_Fisheries_and_Forestry_Management_Using_AI_and_Remote_Sensing
19. [PDF] 14th International Conference “STANDARDIZATION, PROTYPES ... https://www.researchgate.net/profile/Jelena-Ruso-2/publication/328163564_The_Role_of_Higher_Education_in_Quality_Infrastructure_The_Case_of_Serbia/links/61b676631d88475981e6b2f1/The-Role-of-Higher-Education-in-Quality-Infrastructure-The-Case-of-Serbia.pdf
20. (PDF) Bates' Guide to Physical Examination and History Taking https://www.academia.edu/86028490/Bates_Guide_to_Physical_Examination_and_History_Taking
21. 5 - CodaLab Worksheets https://worksheets.codalab.org/rest/bundles/0xd74f36104e7244e8ad99022123e78884/contents/blob/frequent-classes
22. NLP:Bag_of_Words - Kaggle https://www.kaggle.com/code/prashantsparhad/nlp-bag-of-words
23. AI Applications in Game Design | PDF | Artificial Intelligence - Scribd https://www.scribd.com/document/475558209/AInGames
24. [PDF] 14th International Conference “STANDARDIZATION, PROTYPES ... https://www.researchgate.net/profile/Aikaterini-Poustourli-3/publication/328319426_Innovation_Management_building_the_resilient_organisations_of_tomorrow_that_sharing_knowledge/links/5bc7a233458515f7d9c0fe08/Innovation-Management-building-the-resilient-organisations-of-tomorrow-that-sharing-knowledge.pdf

25. Neuroergonomics - The Brain at Work | PDF - Scribd https://www.scribd.com/document/551029961/Neuroergonomics-The-brain-at-work
26. SPAM E-MAILS USING NAIVE BAYES! - Kaggle https://www.kaggle.com/code/hamedetezadi/spam-e-mails-using-naive-bayes
27. Requested packages, organized by age - Debian https://www.debian.org/devel/wnpp/requested_byage
28. Dictionary of Acronyms and Technical Abbreviations - Springer Link https://link.springer.com/content/pdf/10.1007/978-1-4471-3396-4.pdf?qlN92G2paxH3DO=6LzFtLdRT
29. Dictionary of Acronyms and Technical Abbreviations - Springer Nature https://link.springer.com/content/pdf/10.1007/978-1-4471-0263-2.pdf
30. [PDF] HCI in Business, Government and Organizations - ResearchGate https://www.researchgate.net/profile/Muh-Chyun-Tang/publication/372541979_Using_Co-word_Network_Community_Detection_and_LDA_Topic_Modeling_to_Extract_Topics_in_TED_Talks/links/66cbf02575613475fe7b75ff/Using-Co-word-Network-Community-Detection-and-LDA-Topic-Modeling-to-Extract-Topics-in-TED-Talks.pdf
31. Moving technology ethics at the forefront of society, organisations ... https://www.academia.edu/115525006/Moving_technology_ethics_at_the_forefront_of_society_organisations_and_governments
32. ShmooCon dataset (v02) - Greg Conti https://www.gregconti.com/publications/202507-ShmooCon-data-v02.xlsx
33. 6 - CodaLab Worksheets https://worksheets.codalab.org/rest/bundles/0xadf98bb30a99476ab56ebff3e462d4fa/contents/blob/glove.6B.100d.txt-vocab.txt
34. (PDF) `New Game Physics: Added Value for Transdisciplinary ... https://www.academia.edu/5581935/New_Game_Physics_Added_Value_for_Transdisciplinary_Teams_Ph_D_Thesis
35. Spam Message Detector - Kaggle https://www.kaggle.com/code/dev0914sharma/spam-message-detector
36. The African book industry: trends, challenges & opportunities for ... https://unesdoc.unesco.org/ark:/48223/pf0000394048
37. The liminoid-liminal continuum - ResearchGate https://www.researchgate.net/figure/The-liminoid-liminal-continuum_fig1_44911098
38. The Texture of Entrepreneurship Programs: Revisiting Experiential ... https://journals.sagepub.com/doi/10.1177/2515127419890341
39. [PDF] I nternational J ournal of Research in Choral Singing https://acda-publications.s3.us-east-2.amazonaws.com/IJRCS/IJRCSVol10_FullPDF.pdf

40. [PDF] " Who controls the vocabulary, controls the knowledge" - AMiner https://static.aminer.org/pdf/PDF/000/217/452/gsma_a_structural_matching_algorithm_for_schema_matching_in_data.pdf
41. [XLS] China https://www.nsfc.gov.cn/Portals/0/fj/fj20240205_01.xlsx
42. Download book PDF - Springer Nature https://link.springer.com/content/pdf/10.1007/978-3-540-75504-3.pdf
43. PhiloEditor: simplified HTML markup for interpretative pathways over ... https://www.academia.edu/44288549/PhiloEditor_simplified_HTML_markup_for_interpretative_pathways_over_literary_collections
44. Ugc Net Paper 1 | PDF | Educational Technology - Scribd https://www.scribd.com/document/517949987/56000-ugc-net-paper-1
45. suse-sles-15-sp3-sapcal-v20240819-hvm-ssd-x86_64 Package ... https://publiccloudimagechangeinfo.suse.com/amazon/suse-sles-15-sp3-sapcal-v20240819-hvm-ssd-x86_64/package_changelogs.html
46. Dictionary of Acronyms and Technical Abbreviations - Academia.edu https://www.academia.edu/64999604/Dictionary_of_Acronyms_and_Technical_Abbreviations
47. Information and Communications Technology as an Integrity Tool in ... https://www.academia.edu/33311484/Information_and_Communications_Technology_as_an_Integrity_Tool_in_Managing_Climate_change_in_the_Philippines
48. Proceedings of Eighth International Congress on Information and ... https://link.springer.com/content/pdf/10.1007/978-981-99-3043-2.pdf
49. Diksha Paper 1 E Book PDF | PDF | Cognition | Science - Scribd https://www.scribd.com/document/487806939/Diksha-Paper-1-E-Book-pdf
50. Appl. Sci., Volume 10, Issue 23 (December-1 2020) – 452 articles https://www.mdpi.com/2076-3417/10/23
51. ROR in 2025: Year in Review - Research Organization Registry (ROR) https://doi.org/10.71938/xwf9-xd80
52. cmnt_vocab.txt - CMU School of Computer Science https://www.cs.cmu.edu/~ark/blog-data/data/blog_data_v1_0/dk/hbc_data/data/cmnt_vocab.txt
53. (DOC) Stealing fire (word) - Academia.edu https://www.academia.edu/45638035/Stealing_fire_word_
54. Neural Recording Power Optimization Through Machine Learning ... https://www.researchgate.net/publication/396967773_Neural_Recording_Power_Optimization_Through_Machine_Learning_Guided_Resolution_Reconfiguration

55. BioWolf: A sub-10 mW 8-channel Advanced Brain Computer ... https://www.researchgate.net/publication/334352949_BioWolf_A_sub-10_mW_8-channel_Advanced_Brain_Computer_Interface_Platform_with_a_9-core_processor_and_BLE_connectivity
56. BioGAP: a 10-Core FP-capable Ultra-Low Power IoT Processor, with ... https://www.researchgate.net/publication/372709041_BioGAP_a_10-Core_FP-capable_Ultra-Low_Power_IoT_Processor_with_Medical-Grade_AFE_and_BLE_Connectivity_for_Wearable_Biosignal_Processing
57. [DOC] 8203875f2.docx - ResearchGate https://www.researchgate.net/publication/325023504_Supplementary_2/data/5af24259a6fdcc24364f0e63/8203875f2.docx
58. Nefeli Lagoudaki Abused by k Mitsotakis- Kon Samaras-chr Tzermias https://www.scribd.com/document/841434723/Nefeli-Lagoudaki-Abused-by-k-Mitsotakis-Kon-Samaras-chr-Tzermias-finalest
59. Hw3 Stats Google 1gram | PDF | Internet Forum | Software - Scribd https://www.scribd.com/document/917428277/Hw3-Stats-Google-1gram
60. CAFÉ CAVA - Updated March 2026 - 21 Photos & 54 Reviews - Yelp https://www.yelp.com/biz/caf%C3%A9-cava-m%C3%BCnchen
61. La vera vittoria è sconfiggere le discriminazioni❤️ #EqualGame https://www.tiktok.com/@sespo/video/6971470233131388166
62. On of my favourite days out here! #lighthousekeeper ... - TikTok https://www.tiktok.com/@kendra.luckow/video/7560490002010869013
63. (PDF) Designing for Emotions - Academia.edu https://www.academia.edu/34470002/Designing_for_Emotions
64. bing.txt - FTP Directory Listing ftp://ftp.cs.princeton.edu/pub/cs226/autocomplete/bing.txt
65. Live Biofeedback as a User Interface Design Element - ResearchGate https://www.researchgate.net/publication/328033536_Live_Biofeedback_as_a_User_Interface_Design_Element_A_Review_of_the_Literature
66. Abstracts The Israel Society for Neuroscience 18th Annual Meeting https://link.springer.com/article/10.1007/s12031-009-9309-1
67. automated agency and the architecture of survival defining the ... https://www.academia.edu/164810814/AUTOMATED_AGENCY_AND_THE_ARCHITECTURE_OF_SURVIVAL_DEFINING_THE_LUDOLOGICAL_IDENTITY_OF_THE_BULLET_HEAVEN_GENRE_1_
68. High-frequency Broadband Modulations of Electroencephalographic ... https://www.researchgate.net/publication/41029740_High-frequency_Broadband_Modulations_of_Electroencephalographic_Spectra

69. Digital Marketing for Customer Delight | PDF - Scribd https://www.scribd.com/document/893840674/Digital-Marketing-for-Customer-Delight
70. All Audio Logs, Safe Locations & Safe Code Locations https://steamcommunity.com/app/1593230/discussions/0/682985658886660456/
71. Electronics, Volume 14, Issue 19 (October-1 2025) – 222 articles https://www.mdpi.com/2079-9292/14/19
72. Understanding the Brain Function and Emotions - Springer Link https://link.springer.com/content/pdf/10.1007/978-3-030-19591-5.pdf
73. Comprehensive Glossary of Telecom | PDF - Scribd https://fr.scribd.com/doc/266101671/Comprehensive-Glossary-of-Telecom
74. words-333333.txt - cs.Princeton https://www.cs.princeton.edu/courses/archive/fall20/cos226/assignments/autocomplete/files/words-333333.txt
75. Spatial Guidelines for Supernumerary Robotic Limbs in Near-Body ... https://arxiv.org/html/2602.00494v1
76. [PDF] IWSDS 2026 The 16th International Workshop on Spoken Dialogue ... https://aclanthology.org/2026.iwsds-1.pdf
77. (PDF) Library for universal virtual reality experiments (luVRe) https://www.researchgate.net/publication/351862261_Library_for_universal_virtual_reality_experiments_luVRe_A_standardized_immersive_3D360_picture_and_video_database_for_VR_based_research
78. Agentic AI in Healthcare: Autonomous Scheduling and Predictive ... https://www.linkedin.com/posts/shafiqrab_ces-2026-las-vegas-a-few-things-note-worthy-activity-7414475259430584320-OVi6